



**MULTIMEDIA UNIVERSITY**

**OF KENYA**

**FACULTY OF COMPUTING**

**INFORMATION TECHNOLOGY**

**DEPARTMENT OF COMPUTER SCIENCE**

**COURSE: SOFTWARE ENGINEERING**

**Name: Danny Ngatia Gitonga**

**Admission no. : CIT-227-O67/2024**

**Operating Systems – II Review Questions Week**

**3 - Answer all the Questions**

**Remote Procedure Calls**

1. Evaluate the effectiveness of RPCs compared to message-passing systems for implementing distributed applications. Justify your answer based on performance, transparency, and error handling.  
>> ***RPC is more developer-friendly and transparent but less performant and harder to make robust under failure. Message-passing is better for performance-critical or highly reliable custom protocols.***

	<b><u>RPCS</u></b>	<b><u>Message passing systems</u></b>
<b>Performance:</b>	<i>Reduce overhead by adding abstraction layers allowing direct procedure invocation</i>	<i>provide better performance in asynchronous, loosely-coupled scenarios where buffering and queuing are beneficial</i>
<b>Transparency:</b>	<i>Hide network complexity hence invoking procedural calls as if they were local</i>	<i>Exposes messages hence requiring explicit handling of message formats, endpoints and sequencing hence reducing transparency</i>
<b>Error handling:</b>	<i>struggle with network failures, requiring additional mechanisms for handling partial failures, timeouts, and idempotency</i>	<i>handle errors more gracefully through timeouts, acknowledgments, and retry mechanisms built into the messaging infrastructure</i>

2. Given two implementations of an RPC system—one using static binding and the other dynamic binding—which would you recommend for a highly scalable cloud-based service, and why?  
>>***Static binding: Port/endpoint resolved at compile or load time.***  
>>***Dynamic binding: matchmaker/registry (e.g., service discovery via DNS, etcd, or gRPC's name resolver).***  
>>***I would recommend dynamic binding since it offers essential cloud-native scalable services(eg. Kubernetes, microservices) where endpoints change frequently and services auto scale***

3. Critically assess the trade-offs between exactly-once and at-least-once invocation semantics in RPC. Which would you recommend for financial transaction systems, and why?
  - >> **At-least-once:** *Guarantees delivery but may duplicate requests (e.g., due to timeouts). Simpler to implement but risks double-charging.*
  - Exactly-once:** *Ensures a request is processed once and only once. Requires idempotency keys, deduplication logs, or transactional coordination (e.g., two-phase commit).*
  - >> **I would recommend Exactly-once as mandatory for financial transactions because it preserves consistency and correctness, despite higher complexity.**
4. An organization is choosing between implementing gRPC and a custom lightweight RPC framework. As a system analyst, evaluate the two options based on extensibility, maintainability, and support for multiple languages.

	<u><b>gRPC</b></u>	<u><b>Custom lightweight RPC framework</b></u>
<b>Extensibility:</b>	<i>High (interceptors, middleware, streaming)</i>	<i>Limited to in-house design</i>
<b>Maintainability:</b>	<i>Excellent (mature, well-documented, community support)</i>	<i>High long-term cost (testing, bug fixes, docs)</i>
<b>Support for multiple languages:</b>	<i>First-class support (Protobuf + 10+ languages)</i>	<i>Requires manual bindings per language</i>

5. Judge the suitability of RPC in a real-time system (e.g., embedded control system in aviation). What are the critical limitations, and would you recommend using RPC in such a scenario?
  - >> **Critical limitations:**
    - ✓ **Unbounded latency:** *Network delays, retries, or GC pauses violate hard deadlines.*
    - ✓ **Non-determinism:** *RPC relies on OS/network stacks not designed for real-time guarantees.*
    - ✓ **Failure modes:** *Timeouts or lost messages can't be tolerated in safety-critical contexts.*
  - >> **I would not recommend using RPC as it is not suitable for hard real-time. Only acceptable in soft real-time, with caution.**
6. You are tasked with designing a distributed file storage system. Evaluate whether RPC or RESTful APIs would be more appropriate for client-server communication. Justify your choice.
  - >> **RPC (e.g., NFS, gRPC):**
    - *Supports binary protocols, efficient for file ops (read/write blocks).*
    - *Enables stateful sessions, strong consistency, and low-latency batch operations.*
  - >> **REST:**
    - *Stateless, text-based (JSON/XML), higher overhead.*
    - *Better for web clients, caching (HTTP), and loose coupling—but poor for streaming or partial updates.*
  - >> **Recommendation:** *RPC would be more appropriate for client-server communications as it is more appropriate for performance-sensitive file storage (e.g., cloud storage backends like Ceph or HDFS use RPC-like protocols).*
7. Evaluate the fault tolerance mechanisms typically used in RPC systems. Are they sufficient for mission-critical applications? What would you recommend improving?
  - **Mechanisms used are like :** *retries, timeouts, duplicate suppression, client-side caching*  
*They help, but aren't enough for mission-critical apps. Recommend adding checkpointing, replication, and consensus protocols.*

8. Assess the impact of asynchronous RPC on system responsiveness and resource utilization in a microservices architecture. Would you recommend asynchronous over synchronous RPC in such contexts?

>> **Impact:**

- **Responsiveness:** *Clients don't block; better user experience.*
- **Resource utilization:** *Threads/connections freed quickly; higher throughput.*
- **Complexity:** *Requires callbacks, futures, or event loops; harder debugging.*

>> **Yes I would recommend use of asynchronous RPC in microservices to avoid cascading failures and improve scalability (e.g., gRPC async, message queues like Kafka for decoupling) especially where responsiveness matters.**

9. A university student project team wants to use RPC over HTTP to develop a distributed voting system. Critique their approach and suggest whether this is advisable. Support your judgment with reasoning.

>> **RPC over HTTP can be a suitable communication protocol between trusted or internal components of the voting system (e.g., a front-end server communicating with an internal database service)**

**However, it is not advisable for the core voting submission mechanism to the central component without incorporating higher-level, specialized protocols. The core challenges of a distributed voting system are not merely communication efficiency, but securing the fundamental properties of the election itself:**

10. Compare and evaluate the use of middleware frameworks (like CORBA, Java RMI, or gRPC) for implementing RPC in a distributed e-commerce platform. Which framework would you recommend, and on what basis?

Frameworks	pros	cons
<b>COBRA</b>	<i>Language-neutral, mature</i>	<i>Complex, legacy, poor tooling</i>
<b>Java RMI</b>	<i>Simple for Java-only apps</i>	<i>Java-only, no web support</i>
<b>gRPC</b>	<i>High performance, HTTP/2, Protobuf, multi-language, streaming</i>	<i>Steeper learning curve</i>

>> **I would recommend gRPC—best balance of performance, scalability, and polyglot support for modern e-commerce (e.g., order, payment, inventory microservices).**

## **Distributed Processing**

1. Evaluate the design principles of a distributed operating system (DOS). Which principle (e.g., transparency, fault tolerance, scalability) do you consider most critical for system performance, and why?

>> **Fault tolerance is most critical — without it, failures bring down the whole system.**

**Scalability and transparency matter, but reliability is fundamental.**

2. Given a choice between a centralized system and a distributed operating system for a smart city infrastructure project, which would you recommend? Justify your recommendation based on system requirements such as fault tolerance, scalability, and responsiveness.

>> *I would recommend Distributed OS as it is essential for resilience and scale in smart infrastructure, offering:*

- *Fault tolerance: Local failures don't cripple the whole city.*
- *Scalability: Add traffic, energy, or surveillance nodes seamlessly.*
- *Responsiveness: Edge processing reduces latency (e.g., real-time traffic control).*

3. Critique the rationale for adopting distributed systems in large-scale enterprises. Are the benefits (e.g., resource sharing, modular growth) always worth the increased complexity and overhead?

>> *Benefits: resource sharing, modular growth, reliability.*

>> *Downside: higher complexity/overhead.*

>> *Yeah, still worth it for large-scale enterprises where resilience outweighs complexity.*

4. Assess the effectiveness of location transparency in distributed operating systems.

When could this feature become a liability rather than an advantage?

*Advantage: Users/apps don't need to know where data/services reside.*

*Liability:*

- *Performance: Accessing a distant replica increases latency.*
- *Compliance: Data residency laws (e.g., GDPR) require knowing physical location.*
- *Debugging: Obscures root cause of failures.*

*Transparency should be optional—allow apps to opt out for performance or legal reasons.*

5. You are tasked with building a distributed application across a heterogeneous network of devices. Evaluate how the principles of distributed OS design (such as transparency and concurrency) help or hinder your objective.

*Helps:*

- *Transparency abstracts hardware differences (CPU, OS).*
- *Concurrency enables parallel task execution across devices.*

*Hinders:*

- *Heterogeneity complicates data representation (endianness, word size)—requires standard formats (e.g., XDR, Protobuf).*
- *Resource disparity (e.g., IoT vs. server) challenges load balancing.*

6. Judge the appropriateness of using a distributed operating system to manage resources in a university campus network. What limitations or risks should be considered?

>> *Distributed OS helps manage shared resources.*

>> *Risks: admin complexity, higher maintenance costs, and possible security issues.*

7. Evaluate the role of fault tolerance and recovery mechanisms in a distributed OS versus a traditional centralized OS. Are the mechanisms in distributed OSs sufficient for mission-critical systems?

>> *Centralized: Backups and RAID help, but single point of failure remains.*

>> *Distributed: Replication, consensus, and automatic failover provide high availability.*

>> *Modern distributed OSs (e.g., Google Borg, Kubernetes) are sufficient when designed with Byzantine fault tolerance, quorum writes, and chaos engineering.*

8. Compare and assess two architectures for distributed systems: peer-to-peer (P2P) and client-server. Which architecture better supports the principles of a distributed OS and under what conditions?

<u>Aspect</u>	<u>P2P</u>	<u>Client-server</u>
Transparency	<i>Harder (dynamic nodes)</i>	<i>Easier (central naming)</i>
Scalability	<i>Highly scalable</i>	<i>Limited by server</i>
Resource sharing	<i>Organic, efficient</i>	<i>Controlled</i>
Fault tolerance	<i>Resilient (no central node)</i>	<i>Server = SPOF</i>

>> *For distributed OS principles, P2P is usually better unless strong central control is needed.*

9. A startup is considering implementing a distributed system for its logistics operations. Evaluate whether this decision is suitable at their current scale and justify what conditions must be met for distributed processing to be beneficial.

*>> Not suitable initially. Startups need speed, simplicity, and low cost. A monolith or simple cloud app (e.g., AWS Lambda + RDS) suffices.*

*Conditions to adopt distributed processing:*

- *10K daily orders*
- *Need for real-time tracking across regions*
- *SLA requiring 99.9% uptime*
- *Multiple warehouses/suppliers*

10. Assess the trade-offs between performance and transparency in distributed operating systems. Should system designers prioritize one over the other? Defend your position with examples.

*>> Trade-off is unavoidable. For system design, balance is needed, but **performance should take priority** in critical systems (e.g., banking), while transparency can be favored in user-facing systems.*